

Machine Learning Notes

Theory, Practice, and Interview Review

Erdan Mu Alan

University of Pennsylvania

Contents

I	Foundations	1
II	Statistical Learning	2
1	Linear Regression	3
2	Logistic Regression	4
3	Model Evaluation	5
III	Classical Machine Learning	6
4	Distance-Based Learning	7
5	Support Vector Machines	8
6	Decision Trees	9
7	Ensemble Learning	10
IV	Unsupervised Learning	11
8	Clustering	12
9	Dimensionality Reduction	13
10	Density Estimation	14
V	Deep Learning Fundamentals	15
11	Neural Networks	16
12	Optimization in Deep Learning	17
13	Regularization in Deep Learning	18
VI	Deep Learning Architectures	19
14	Sequence Models	20

CONTENTS	ii
15 Attention and Transformers	21
16 Graph Neural Networks	22
VII Representation Learning and Generative Models	23
17 Representation Learning	24
18 Generative Models	25
VIII Large Language Models	26
19 LLM Foundations	27
20 Post-Training	28
21 Retrieval and Agents	29
IX Practical Machine Learning	30
22 Data Processing	31
22.1 Missing Values	31
22.2 Feature Engineering	31
22.3 Encoding	31
22.4 Scaling	31
22.5 Data Leakage	31
23 Experimentation	32
23.1 Hyperparameter Tuning	32
23.2 Grid Search	32
23.3 Random Search	32
23.4 Bayesian Optimization	32
23.5 Ablation Studies	32
24 MLOps and Deployment	33
24.1 Model Serving	33
24.2 Monitoring	33
24.3 Drift Detection	33
24.4 Retraining Pipelines	33
X Interview Questions	34

Part I

Foundations

Part II

Statistical Learning

Chapter 1

Linear Regression

Can anyone tell me what happened?

Chapter 2

Logistic Regression

Chapter 3

Model Evaluation

Part III

Classical Machine Learning

Chapter 4

Distance-Based Learning

Chapter 5

Support Vector Machines

Chapter 6

Decision Trees

Chapter 7

Ensemble Learning

Part IV

Unsupervised Learning

Chapter 8

Clustering

Chapter 9

Dimensionality Reduction

Chapter 10

Density Estimation

Part V

Deep Learning Fundamentals

Chapter 11

Neural Networks

Chapter 12

Optimization in Deep Learning

Chapter 13

Regularization in Deep Learning

Part VI

Deep Learning Architectures

Chapter 14

Sequence Models

Chapter 15

Attention and Transformers

Chapter 16

Graph Neural Networks

Part VII

Representation Learning and Generative Models

Chapter 17

Representation Learning

Chapter 18

Generative Models

Part VIII

Large Language Models

Chapter 19

LLM Foundations

Chapter 20

Post-Training

Chapter 21

Retrieval and Agents

Part IX

Practical Machine Learning

Chapter 22

Data Processing

22.1 Missing Values

22.2 Feature Engineering

22.3 Encoding

22.4 Scaling

22.5 Data Leakage

Chapter 23

Experimentation

23.1 Hyperparameter Tuning

23.2 Grid Search

23.3 Random Search

23.4 Bayesian Optimization

23.5 Ablation Studies

Chapter 24

MLOps and Deployment

24.1 Model Serving

24.2 Monitoring

24.3 Drift Detection

24.4 Retraining Pipelines

Part X

Interview Questions